Prompting LLMs Using Haystack

Experiment with PromptNode and new prompt templates

Haystack Webinar | Vladimir Blagojevic | April 4th, 2023

Agenda

1	Introduction
2	Prompt Engineering - Fundamentals & Guidelines
3	Common Prompting Pitfalls
4	Latest Research
5	Q&A

01 Introduction

LLMs and their Impact on Our Society

- “The hottest new programming language is English” - Andrej Karpathy
- LLMs represent a world model
- Innovative application and transformative potential
- Prompting is a critical skill to harness LLMs
- Haystack LLMs API abstractions: `PromptNode`, `PromptTemplate`

Recap: What's available in Haystack?

`PromptNode`

A node that allows you to interact with LLMs in a customized way. It allows you to use models from OpenAI and Hugging Face.

`PromptTemplate`

A component that allows you to build flexible prompts that can be modified per query. This component can be provided To the `PromptNode` as a blueprint of how to interact with LLMs

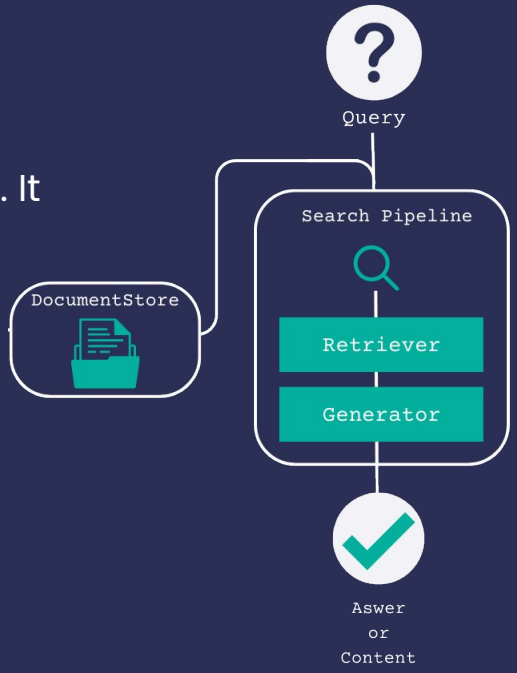

02

Prompt Engineering – Fundamentals & Guidelines

Fundamentals of Prompt Engineering

Why fundamentals?

Context window limitations

Attention mechanism

Model name	Window
flan-t5	512
text-davinci-003	4,097
gpt-4	8,192 32,768

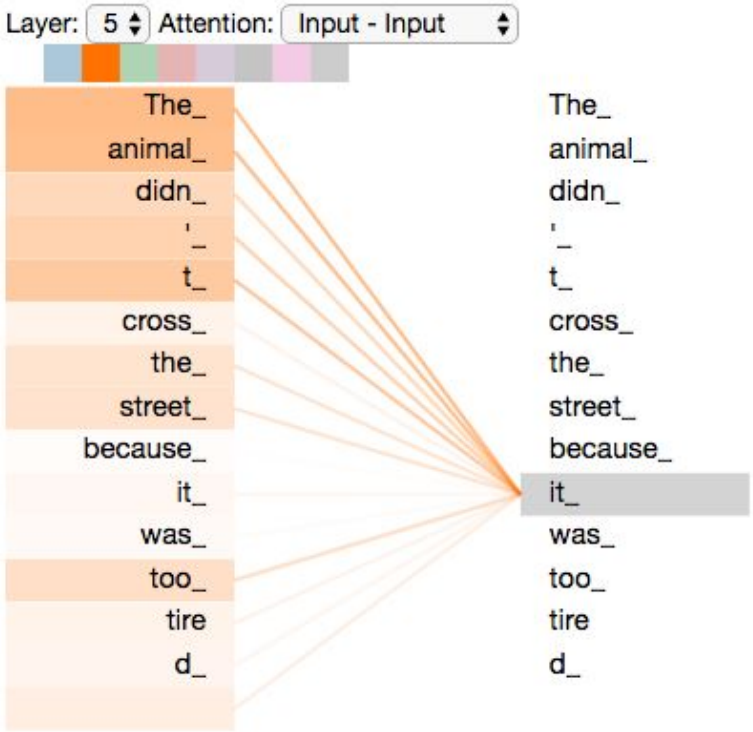

Fundamentals of Prompt Engineering

```
[28] pn = PromptNode("text-davinci-003", api_key=openai_key)
     poem = pn("Write a short poem about Haystack")
     print(poem[0])
```

```
Rock

A rock of grandeur stands tall in the sea,
A landmark admired by all, one can see.
It stands as a symbol of nature's might,
A beacon of beauty, Haystack Rock's light.
```

```
[29] prompt = """
     "Write a short poem about Haystack,
     an open-source NLP framework,
     focus on recently released Haystack PromptNode and its powerful templates"""
     poem = pn(prompt)
     print(poem[0])
```

```
Haystack PromptNode is here
It brings new tools and templates without fear
Enabling NLP task automation with ease
Haystack's power is clear!
```

Clarity and specificity

Fundamentals of Prompt Engineering

Incorporating leading words (e.g. “let’s think step by step”)

```
[12] pn = PromptNode("text-davinci-003", api_key=openai_key)
```

0s

```
[13] prompt = """
     "Add 1 and 3 and then divide by the latter and then add the former.
     Don't say anything else other than the answer.
     Not one word."""
     pn(prompt)
```

```
['2']
```

0s

```
[14] prompt = """
     Add 1 and 3 and then divide by the latter and then add the former.
     Let's think step-by-step:
     """
     pn(prompt)
```

```
['1. Add 1 and 3: 1 + 3 = 4 \n2. Divide by 3: 4/3 = 1.33 \n3. Add 1: 1.33 + 1 = 2.33']
```

1s

Fundamentals of Prompt Engineering

Yeah, ok, but GPT-4 doesn't need leading words, right?!

Model: GPT-4

Add 1 and 3 and then divide by the latter and then add the former. Don't say anything else other than the answer. Not one word.

2

Fundamentals of Prompt Engineering

```
[31] prompt_text = """Extract keywords from the below text.
     Text: {text}
     Keywords:"""

     text = """
     Former President Donald Trump was initially asked to turn
     himself in to authorities in New York today — the day
     after a Manhattan grand jury voted to indict him, his
     defense lawyer Joe Tacopina says.
     """
     pt = PromptTemplate("extract-keywords-zero-shot", prompt_text)
     pn.prompt(prompt_template=pt, text=text)

     ['Former President Donald Trump, turn himself in, authorities, New York, Manhattan, grand jury, indict,
```

Start zero-shot

Fundamentals of Prompt Engineering

Continue with few-shot

```
[32] prompt_text = """
     Extract keywords from the corresponding texts below.

     Text 1: Stripe provides APIs that web developers can use to integrate payment processing into
     Keywords 1: Stripe, payment processing, APIs, web developers, websites, mobile applications
     ##
     Text 2: OpenAI has trained cutting-edge language models that are very good at understanding an
     Keywords 2: OpenAI, language models, text processing, API.
     ##
     Text 3: {text}
     Keywords 3:
     """
     pt = PromptTemplate("extract-keywords-few-shot", prompt_text)
     pn.prompt(prompt_template=pt, text=text)

     ['Donald Trump, authorities, New York, Manhattan grand jury, Joe Tacopina.']
```

Prompt Engineering Guidelines

- Provide clear and explicit instructions
- Include examples if needed
- Experiment with different prompt styles
- Utilize system-level constraints (temperature, token limits)
- Iterate and refine prompts based on feedback
- Address potential biases in LLMs and prompts

Prompt Engineering Guidelines

```
[ ]  prompt_text = """
     Describe {query} in a few sentences only. Use documents provided below.
     Documents: {join(documents)}
     """

     prompt_text = """
     Please read the text provided:

     {join(documents)}

     Craft a 3-5 sentence description about {query}, ensuring it is concise and informative.
     Start by introducing the ski's name and type, then discuss the ski building technology
     used, and finally highlight the main benefits.
     Use skier-specific terminology for a more engaging and relatable description, drawing
     information exclusively from the provided text.
     """


     prompt_node = PromptNode(
         "text-davinci-003",
         default_prompt_template=PromptTemplate("product_description", prompt_text=prompt_text),
         api_key=openai_key,
         max_length=256,
     )

     web_retriever = WebRetriever(api_key=search_key, top_search_results=5, mode="preprocessed_documents")
     pipeline = WebQAPipeline(retriever=web_retriever, prompt_node=prompt_node)
     output = pipeline.run("Dynastar speed 963")
```

Summary of everything we learned

Prompt Engineering Guidelines

```
prompt_text = """
Describe {query} in a few sentences only. Use documents provided below.
Documents: {join(documents)}
"""
```

```
[53] output["results"][0]
```

```
'The traditional sandwich construction gives the ski precise flex and rebound control.\n\nThe Dynastar Speed 963 is a top-of-the-line alpine ski that combines performance and power with lightweight agility. Featuring a hybrid core for a unique on-snow feeling, a Poplar Bi Directional for better resistance to compression, V Tech for enhanced rigidity and torsional control, Tip Rocker for a softer and more tolerant curve, and a Sandwich Construction for precise flex and rebound control, the Speed 963 offers an unparalleled skiing experience.'
```

```
prompt_text = """
Please read the text provided:

{join(documents)}

Craft a 3-5 sentence description about {query}, ensuring it is concise and informative.
Start by introducing the ski's name and type, then discuss the ski building technology
used, and finally highlight the main benefits.
Use skier-specific terminology for a more engaging and relatable description, drawing
information exclusively from the provided text.
"""
```

```
[51] output["results"][0]
```

```
'The Dynastar Speed 963 is a perfect choice for skiers looking for a reliable performance. It features a Hybrid Core which blends the performance of wood with the lightness and smooth snow feel of PU, Poplar Bi Directional for resistance to compression and dynamic reaction in flexion, V-Tech for rigidity, power and torsional control, and Tip Rocker for facilitated flat pivot and softness in curves. These features provide an exceptional on-mountain feeling, perfect for advanced and expert skiers looking for a reliable, comfortable and agile ride.'
```

Summary of everything we learned

03 Common Prompting Pitfalls

Common Prompting Pitfalls

- Ambiguous and vague instructions
- Overly complex or lengthy prompts
- Single-token answers

04

Latest Research

Prompting and latest research

Automatic Prompt Engineer (APE)

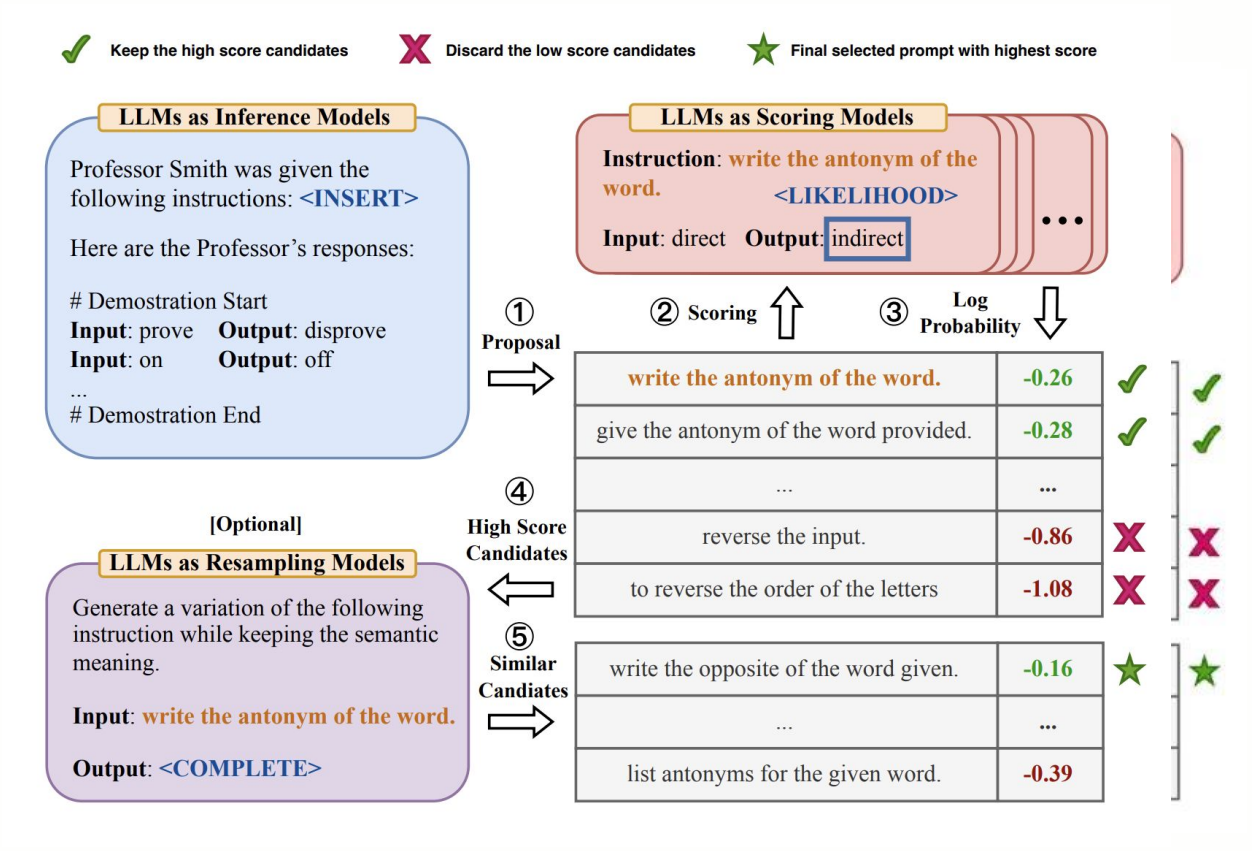

Zhou et al, University of Toronto

Prompting and latest research

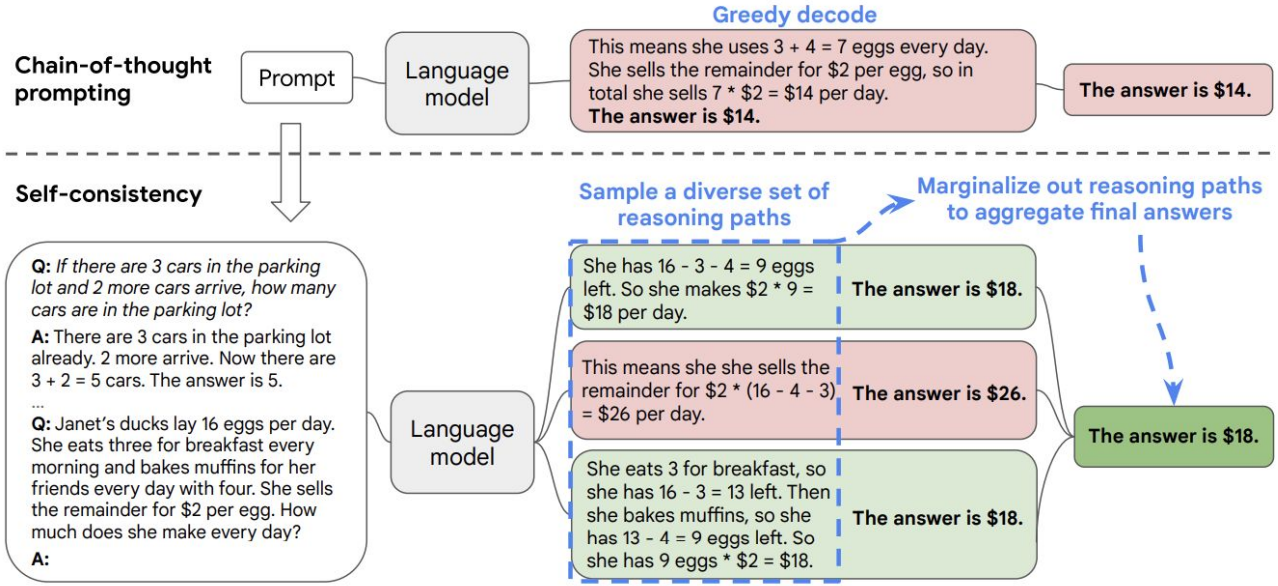

Wang et al. Google Brain

Self-consistency sampling

Prompting and latest research

Self-Taught Reasoner (STaR) method

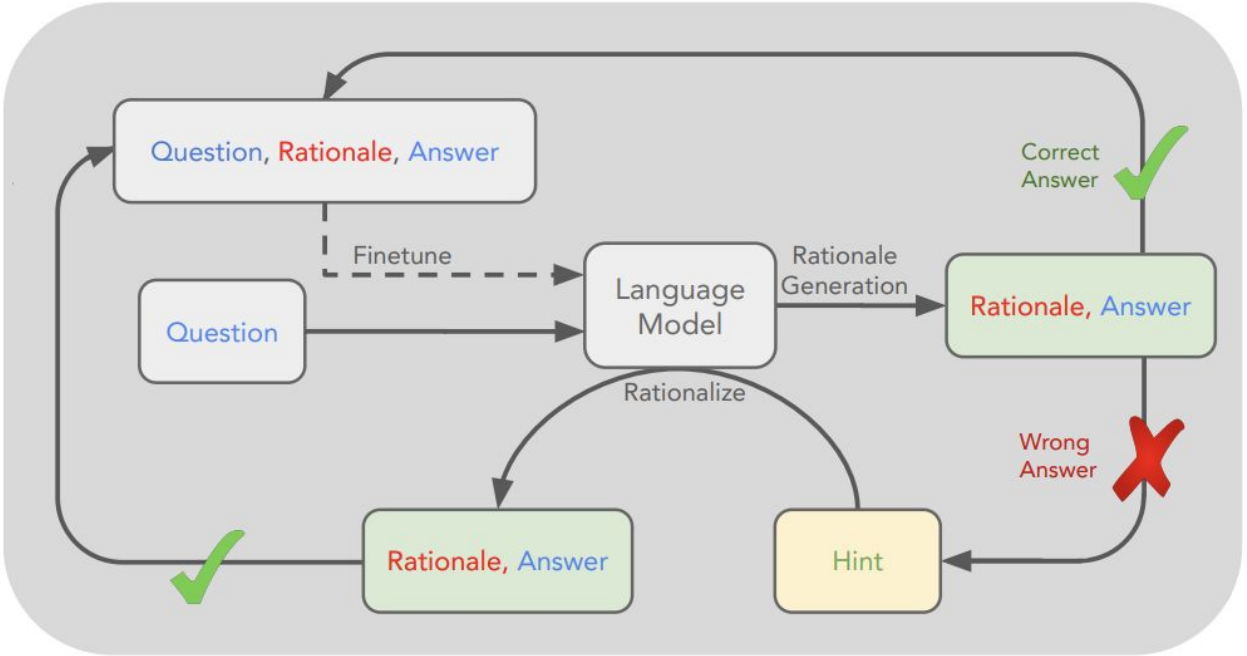

Zelikman et al, Stanford + Google Research

Secret Encore

The ultimate GPT-4 prompter for you to use

Let’s review it quickly

https://tinyurl.com/HaystackPrompt

05

Q&A

List of resources to write down while we discuss

https://www.promptingguide.ai/

https://learnprompting.org/

https://www.getsphere.com/cohorts/prompt-engineering-for-llms

https://lilianweng.github.io/posts/2023-03-15-prompt-engineering/

Thank you!

www.haystack.deepset.ai

Vladimir Blagojevic vladimir.blagojevic@deepset.ai